

# In-Class Assignment 3

```
In [8]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import statsmodels.api as sm
from statsmodels.formula.api import ols
from sklearn.model_selection import train_test_split
import scipy.stats as ss

advertising_df = pd.read_csv("Advertising.csv", index_col=0)
advertising_df['intercept'] = 1
y_var = 'sales'
advertising_df.head()
```

```
Out[8]:   TV  radio  newspaper  sales  intercept
1  230.1    37.8       69.2   22.1          1
2   44.5    39.3       45.1   10.4          1
3   17.2    45.9       69.3    9.3          1
4  151.5    41.3       58.5   18.5          1
5  180.8    10.8       58.4   12.9          1
```

## Estimating the Coefficients

```
In [10]: # x_vars = ['intercept', 'TV']
# model = sm.OLS(advertising_df[y_var], advertising_df[x_vars])

model = ols(f"{y_var} ~ TV", data=advertising_df)
results = model.fit()
results.summary()
```

Out[10]:

## OLS Regression Results

<b>Dep. Variable:</b>	sales	<b>R-squared:</b>	0.612			
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.610			
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	312.1			
<b>Date:</b>	Wed, 04 Sep 2024	<b>Prob (F-statistic):</b>	1.47e-42			
<b>Time:</b>	19:50:55	<b>Log-Likelihood:</b>	-519.05			
<b>No. Observations:</b>	200	<b>AIC:</b>	1042.			
<b>Df Residuals:</b>	198	<b>BIC:</b>	1049.			
<b>Df Model:</b>	1					
<b>Covariance Type:</b>	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	7.0326	0.458	15.360	0.000	6.130	7.935
<b>TV</b>	0.0475	0.003	17.668	0.000	0.042	0.053
<b>Omnibus:</b>	0.531	<b>Durbin-Watson:</b>	1.935			
<b>Prob(Omnibus):</b>	0.767	<b>Jarque-Bera (JB):</b>	0.669			
<b>Skew:</b>	-0.089	<b>Prob(JB):</b>	0.716			
<b>Kurtosis:</b>	2.779	<b>Cond. No.</b>	338.			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

We've optimized the parameters  $\beta_0$  and  $\beta_1$ , what is the objective?

Let  $e_i = y_i - \hat{y}_i$  represent the residual of a prediction  $\hat{y}_i$ . Then the sum of squared residuals is  $RSS = \sum_{i=1}^n e_i^2 = \sum_{i=1}^n (y_i - \hat{y}_i)^2$

We can think of this as the amount of variability that is left in the response after performing the regression.

```
In [12]: RSS = ((results.predict(advertising_df[['intercept', 'TV']]) - advertising_df['sale'])**2).sum()
```

Out[12]: 2102.530583131351

```
In [14]: sm.stats.anova_lm(results)
```

Out[14]:		df	sum_sq	mean_sq	F	PR(>F)
	<b>TV</b>	1.0	3314.618167	3314.618167	312.144994	1.467390e-42
	<b>Residual</b>	198.0	2102.530583	10.618841	NaN	NaN

**Where does the standard error come from? Do we know the true standard deviation of the errors in the model?**

$$\hat{\sigma} = RSE = \sqrt{RSS/(n - p - 1)}$$

```
In [16]: # degrees of freedom = n - p - 1, where p is the number of predictor variables
degrees_of_freedom = (len(advertising_df) - 1 - 1)

RSE = np.sqrt(RSS / degrees_of_freedom)
RSE
```

Out[16]: 3.258656368650462

```
In [6]: SE_beta_1 = np.sqrt((RSE**2) / ((advertising_df['TV'] - advertising_df['TV'].mean())
SE_beta_1.round(3))
```

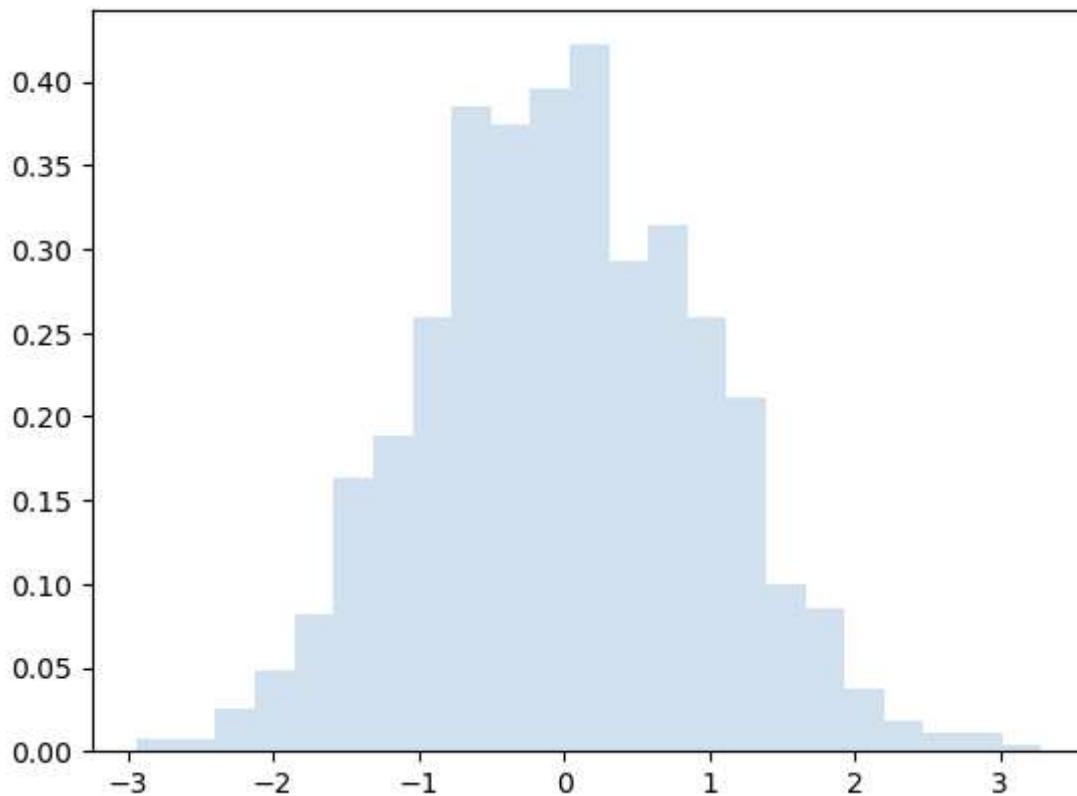
Out[6]: 0.003

## What is this $t$ ?

```
In [7]: t_var = ss.t(degrees_of_freedom)
sample = t_var.rvs(size=1000)

fig, ax = plt.subplots()
ax.hist(sample, density=True, bins='auto', histtype='stepfilled', alpha=0.2)
ax.legend(loc='best', frameon=False)
plt.show()
```

No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.



## Let's calculate a t-statistic, p-value, and confidence interval

- Note: the t-statistic for a parameter is measured with the hypothesis that the parameter equals 0, i.e. it has no relationship with the outcome

```
In [8]: results.params
```

```
Out[8]: Intercept    7.032594
         TV          0.047537
         dtype: float64
```

```
In [9]: beta_1 = results.params['TV']
t_stat_beta_1 = (beta_1 - 0) / SE_beta_1

p_val_t_stat_beta_1 = 1 - t_var.cdf(t_stat_beta_1)

lb_beta_1 = beta_1 + t_var.ppf(.025) * SE_beta_1
ub_beta_1 = beta_1 + t_var.ppf(.975) * SE_beta_1

print("t-statistic:", round(t_stat_beta_1, 3))
print("p-val of t-statistic:", p_val_t_stat_beta_1)
print("95% Confidence Interval:", round(lb_beta_1, 3), "to", round(ub_beta_1, 3))
```

```
t-statistic: 17.668
p-val of t-statistic: 0.0
95% Confidence Interval: 0.042 to 0.053
```

## What proportion of the variance in sales is explained by this model?

Let the sum of squared difference between the outcomes ( $y_i$ ) and the mean of all outcomes ( $\bar{y}$ ) be TSS, i.e.

$$TSS = \sum_{i=1}^n (y_i - \bar{y})^2 \quad (1)$$

$$R^2 = \frac{TSS - RSS}{TSS} \quad (2)$$

```
In [10]: TSS = ((advertising_df['sales'] - advertising_df['sales'].mean())**2).sum()
R_squared = (TSS - RSS) / TSS
R_squared
```

```
Out[10]: 0.6118750508500711
```

## Multiple regression

```
In [37]: model_multiple = ols(f"y_var ~ TV + radio + newspaper", data=advertising_df)
results_multiple = model_multiple.fit()
results_multiple.summary()
```

Out[37]:

## OLS Regression Results

<b>Dep. Variable:</b>	sales	<b>R-squared:</b>	0.897			
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.896			
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	570.3			
<b>Date:</b>	Wed, 04 Sep 2024	<b>Prob (F-statistic):</b>	1.58e-96			
<b>Time:</b>	12:10:53	<b>Log-Likelihood:</b>	-386.18			
<b>No. Observations:</b>	200	<b>AIC:</b>	780.4			
<b>Df Residuals:</b>	196	<b>BIC:</b>	793.6			
<b>Df Model:</b>	3					
<b>Covariance Type:</b>	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	2.9389	0.312	9.422	0.000	2.324	3.554
<b>TV</b>	0.0458	0.001	32.809	0.000	0.043	0.049
<b>radio</b>	0.1885	0.009	21.893	0.000	0.172	0.206
<b>newspaper</b>	-0.0010	0.006	-0.177	0.860	-0.013	0.011
<b>Omnibus:</b>	60.414	<b>Durbin-Watson:</b>	2.084			
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	151.241			
<b>Skew:</b>	-1.327			<b>Prob(JB):</b>	1.44e-33	
<b>Kurtosis:</b>	6.332			<b>Cond. No.</b>	454.	

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

**We can see all of the previous statistics, now let's perform a hypothesis test on whether the model variables have a relationship with the response**

- Another way of stating this is that our hypothesis is that all of the estimated model parameters on the predictor variables are zero, i.e.  $\beta_1 = \beta_2 = \beta_3 = 0$
- We might want to think that we can just look at the p-values of the variables, but this is not valid. If we had 100 variables, we expect about 5 of those to have a p-value of less than .05!

The  $F$  statistic can help us with this:  $F = \frac{(TSS-RSS)/p}{RSS/(n-p-1)}$

- If our model assumptions are correct we expect that the numerator and denominator both equal the variance of the error distribution, i.e.  $F = \frac{\sigma^2}{\sigma^2} = \frac{Var(\epsilon)}{Var(\epsilon)}$  where the model is  $Y \sim f(X) + \epsilon$ .

```
In [38]: TSS_multiple = results_multiple.centered_tss
RSS_multiple = results_multiple.ssr

p = 3
n = len(advertising_df)

F_stat_multiple = ((TSS_multiple - RSS_multiple) / p) / (RSS_multiple / (n - p - 1))

print("TSS:", TSS_multiple)
print("RSS:", RSS_multiple)
print("F-stat:", F_stat_multiple)
```

TSS: 5417.14875  
 RSS: 556.8252629021874  
 F-stat: 570.270703659094

## Questions

- Base all of your answers on the additive model  

$$Y \sim \beta_0 + \beta_1 * TV + \beta_2 * radio + \beta_3 * newspaper.$$
1. Is there a relationship between sales and advertising budget? What test would you use to determine this?
  - Answer: Yes, F test
  2. How strong is the relationship? What statistic would you use to measure this?
    - Answer: 89.7% of the variance is explained by our model based on the  $R^2$  statistic
  3. Which of the media are associated with sales? What statistic would you use to measure this?
    - Answer: TV and Radio, by using t-statistic
  4. For each unit increase in TV advertising budget, what change in sales should we expect? What is a 95% confidence interval of that effect?
    - Answer: We expect an average change of 0.046 in sales for a unit increase in TV ad expenditure. The confidence interval at 95% is [0.043, 0.049]

In [ ]: